

# An Implementation of the Homa Transport Protocol in RAMCloud

Yilong Li, Behnam Montazeri, John Ousterhout



PLATFORMLAB

# Introduction

---

- **Homa: receiver-driven low-latency transport protocol using network priorities**
  - Key ideas and simulation results presented before
- **HomaTransport in RAMCloud: a working implementation**
  - Unusual features: message-oriented, connectionless, no ACKs, etc.
- **Excellent performance**
  - Extreme network condition: **80%** network load on 10 Gbps network
  - Slowdown of 99%-tile latency of almost all message sizes within **2-3.5x**
  - 99%-tile round-trip latencies for small messages < **15**  $\mu$ s
  - Nearly **100x** faster than best published result


# Outline

---

- **Homa Overview**
- **Implementation Features**
- **Evaluation**

# Homa Overview

- **Goal: low latency at high network load**
  - Focus on tail (e.g., 99th percentile) message latency
  - Implement shortest-remaining-processing-time (SRPT) policy
- **Key Idea 1: Divide outgoing messages into unscheduled and scheduled portions**




# Homa Overview

- **Goal: low latency at high network load**
  - Focus on tail (e.g., 99th percentile) message latency
  - Implement shortest-remaining-processing-time (SRPT) policy
- **Key Idea 2: Dynamic priority allocation**
  - Receiver can change the priorities granted to incoming messages on the fly
  - ... based on the exact set of incoming messages



# Homa Overview

- **Goal: low latency at high network load**
  - Focus on tail (e.g., 99th percentile) message latency
  - Implement shortest-remaining-processing-time (SRPT) policy
- **Key Idea 3: Controlled overcommitment**
  - Receiver grants to a smaller number of senders simultaneously
  - ... to avoid wasting downlink bandwidth
  - ... when our most favored sender doesn't send back granted data in time



# Implementation in RAMCloud

---

- **RAMCloud: low-latency key-value store**
  - Modular transport architecture
  - Optimized software stack: 1~2  $\mu$ s to send/receive an RPC
- **RAMCloud::HomaTransport**
  - Kernel bypass via DPDK
  - Use polling to detect incoming packets
  - ~ 4000 lines of C++ code (including comments)

# Homa: Structurally Different From TCP

---

- **Message-oriented, not stream-oriented**
  - Independent delivery of messages: no head-of-line-blocking
  - RPC interface
    - Natural fit for datacenter applications
    - Socket-like byte stream interface on top of Homa
- **Connectionless**
  - No setup phase required before sending an RPC
  - No state kept after RPC completes




# Retransmission

---

- **No explicit ACKs**
  - RPC response as the acknowledgment for the request
  - Optimize for small RPCs: reduce half of the packets required
    - The simplest RPC requires only 1 **DATA** packet in each direction
- **Receiver-driven approach to detect lost packets**
  - Receiver timeouts on messages that have been silent for a long period
    - ... and request retransmission for the missing bytes
  - What if all unscheduled packets of a request are lost?
    - Client eventually timeouts on the response message
    - ... and request retransmission of the initial bytes of the response
    - Server doesn't recognize this RPC, assumes that the request must be lost
    - ... and request retransmission of the initial bytes of the request


# Sender-Side Queue Limiting

- **Sender implements SRPT by default**
  - Need to preempt long messages for short ones
  - Keep the transmit queue in NIC short to avoid queueing delay
- **QueueEstimator**
  - Keep a running estimate of the transmit queue length
- **Limiting queue length**
  - Too large: increase queueing delay for short messages
  - Too small: risk of TX queue running dry
  - Enqueue a packet only if  $\text{queueLength} \leq \text{one full-size packet}$




# Evaluation Experiment

- 8 clients and 8 servers
- Each client generates a series of echo RPCs to random servers
  - Client sends a message of a given size
  - Server replies with the same message
- RPC message size chosen randomly to match the given workload
- RPC inter-arrival times follow poisson distribution
  - Average inter-arrival time configured to generate a given network load



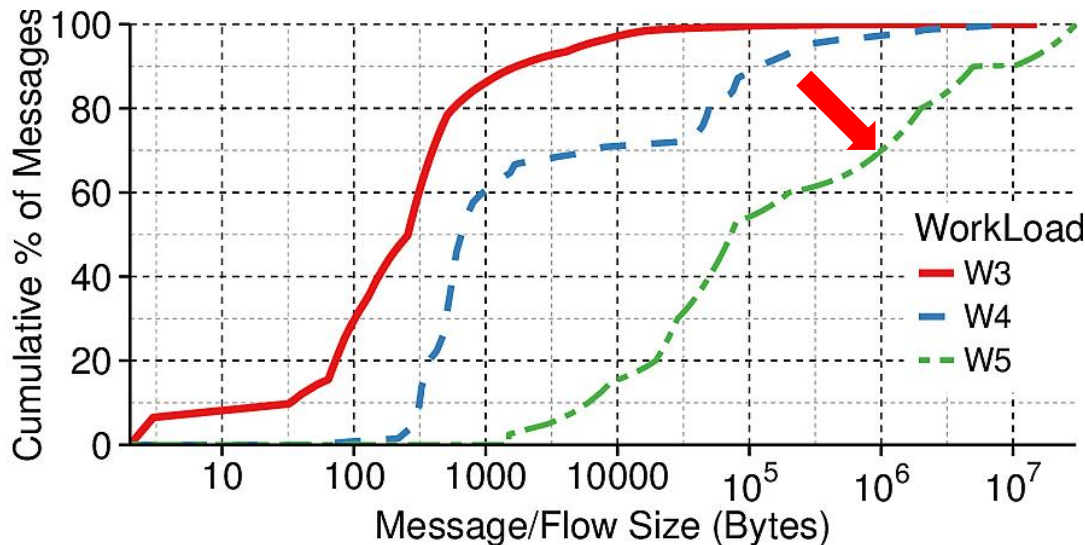
# Workloads

- **Workload: distribution of message sizes**
  - W3: aggregated RPC workload from Google datacenter applications
  - W4: Facebook Hadoop workload
  - W5: web search workload used for DCTCP



# Workloads

- **Workload: distribution of message sizes**
  - W3: aggregated RPC workload from Google datacenter applications
  - W4: Facebook Hadoop workload
  - W5: web search workload used for DCTCP



# Configuration


---

- **Hardware configurations**

- CloudLab m510 cluster: 8-Core Xeon D1548 @ 2.0 GHz, 10 Gbps network
- Local Infiniband cluster: 4-Core Xeon X3470 @ 2.93 GHz, 24 Gbps network
- All nodes in a cluster are connected to a single switch


# Understanding The Graph

Measurements are taken on 10 Gbps network at **80%** network load unless stated otherwise.



# Understanding The Graph

Measurements are taken on 10 Gbps network at **80%** network load unless stated otherwise.






# Understanding The Graph


Measurements are taken on 10 Gbps network at **80%** network load unless stated otherwise.



**Slowdown =**  
**Actual RPC Completion Time /**  
**Best-case RPC Completion Time**  
**(On Unloaded Network)**




# Homa Absolute Performance



## Best-case RPC time

- 100 bytes: 4.7  $\mu$ s
- 1000 bytes: 8.8  $\mu$ s


# Homa Absolute Performance



## Best-case RPC time


- 100 bytes: 4.7  $\mu$ s
- 1000 bytes: 8.8  $\mu$ s

# HOL Blocking = 100x Tail Latency




**Note: InfRC measurements are taken on a 24 Gbps Infiniband network using the absolute same workload, so the actual network load is only 33%.**


# HOL Blocking = 100x Tail Latency



# HOL Blocking = 100x Tail Latency




# Homa vs. TCP



# Priority & Controlled Overcommitment


Basic vs. Homa: 5 - 15x higher tail latency for most RPCs






# Priority & Controlled Overcommitment

Basic vs. Homa: 5 - 15x higher tail latency for most RPCs




# Priority & Controlled Overcommitment

Basic vs. Homa: 5 - 15x higher tail latency for most RPCs



**SRPT tends to produce run-to-completion behavior**


# How Many Priorities Does Homa Need?



**4 priority levels is almost as good as 8.**

# Results Overview

- Homa vs. Other RAMCloud transports at 80% load (except InfRC)



# Conclusion

---

- **We designed and implemented Homa, a new transport protocol for datacenter networks**
  - Provide very low latency for short messages
  - Support high network utilization
- **Our Homa implementation has several unusual features**
  - Message-oriented, connectionless, no explicit ACKs, etc.
- **Implementation measurements show excellent performance numbers across various workloads**

---

# Questions?